

Hardened, scope-locked first-price search (44 valid candidates $\times$ 5-mouse panel, GPT-5.4-mini)

Scope-locked first-price ranking
(grid/scaffolded variants lead)

Efficiency is uniformly high;
the worst-persona tail drives the score

